

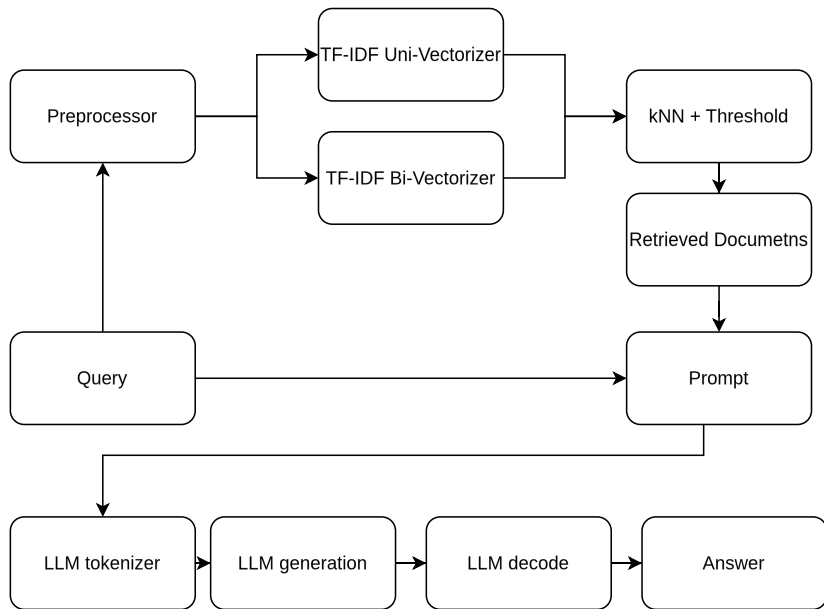
IRSE Project

Jan Cichomski (r1026448)

May 15, 2025

- why have i chose F1 as a metric?

1. Architecture



2. Term Vocabulary

2.1 Term Vocabulary - Document Preprocessing

- To lower case
- Remove punctuation
- Tokenize
- Remove english stop words (added custom stop words)
- Lemmatize

2.1 Temr Vocabulary - Document Preprocessing

```
def preprocess_text(doc):  
    doc = doc.translate(str.maketrans("", "",  
        string.punctuation)).lower()  
    words = word_tokenize(doc)  
    words = [  
        lemmatizer.lemmatize(word)  
        for word in words  
        if word not in stop_words and word.isalpha()  
    ]  
    return " ".join(words)
```

2.1 Term Vocabulary - Custom stop words

```
stop_words.update(  
    [  
        "add",  
        "added",  
        "adding",  
        "addition",  
        "also",  
        "almost",  
        "another",  
        "easily",  
        "easy",  
    ]  
)
```


2.2 Term Vocabulary - Hyperparameters

Two types of terms:

- 1-grams
 - min_df=20
 - max_df=0.5
- 2-grams
 - 10,000 terms
 - min_df=50
 - max_df=0.4

2.3 Term Vocabulary - Handling multi-word terms

- 2-grams with aggressive filtering

3 Document Embedding

3.1 Document Embedding - Chosen Fields

Fields used for embedding:

- name
- description
- ingredients
- steps

Evaluated combinations:

- name, description, ingredients, steps: macro F1: 0.126
- description, ingredients, steps: macro F1: 0.095
- description, ingredients: macro F1: 0.034
- description, steps: macro F1: 0.088
- description: macro F1: 0.043

3.2 Document Embedding - Query Preprocessing

The same approach as for embedding documents

```
def retrieve_documents(query_text, recipes,  
                      recipe_ids, k, threshold):  
    query = preprocess_text(query_text)  
    ...
```

3.3 Document Embedding - Edge Cases

- When query has no terms from vocabulary
 - TF-IDF produces zero vector for the query
 - Cosine similarity returns 0 for all documents
- Consequently:
 - Without similarity threshold: All documents returned (no filtering)
 - With any similarity threshold: No documents returned (empty result)

4 Retrieval

4.1 Retrieval - Similarity Measure

- Cosine similarity (F1: 0.126) - picked finally
- Euclidean distance (F1: 0.064)

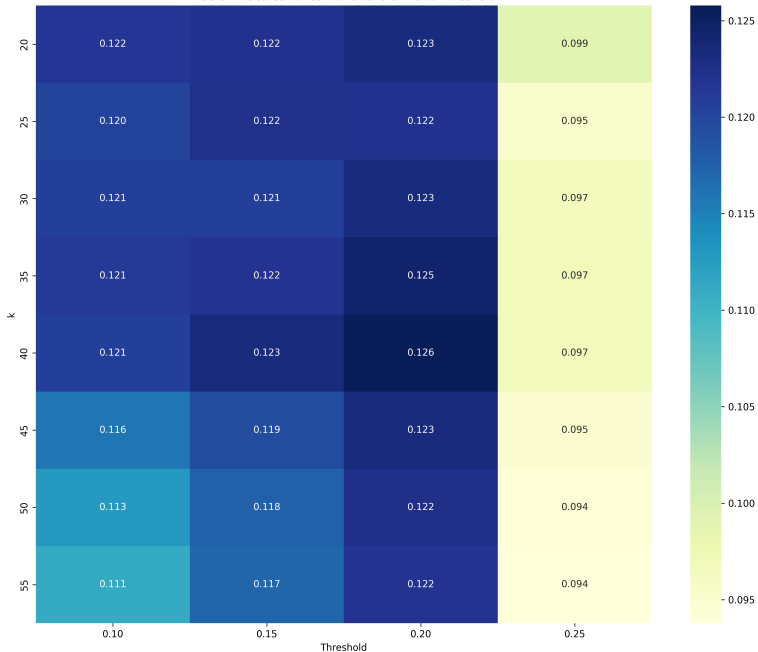
4.2 Retrieval - Hyperparameters

- Max number of returned documents: 40
- Minimum threshold for cosine similarity: 0.2

Grid search over param space and maximize macro F1 score.

```
def create_parameter_heatmap(queries, recipes, recipe_ids):  
    thresholds = np.arange(0.1, 0.60, 0.05)  
    k_values = np.arange(20, 60, 5)
```

Macro F1 Scores for Combinations of k and Threshold



4.3 Retrieval - Evaluation Metrics

- Macro Precision: 0.130
- Macro Recall: 0.201
- Macro F1: 0.126
- Micro Precision: 0.128
- Micro Recall: 0.191
- Micro F1: 0.153

4.4 Retrieval - MAP

Mean Average Precision (MAP): 0.086

$$AP = \frac{1}{RD} \sum_{k=1}^n P(k) \cdot r(k), \quad (1)$$

Where RD is the number of relevant documents for the query, n is the total number of documents, $P(k)$ is the precision at k , and $r(k)$ is the relevance of the k^{th} retrieved document (0 if not relevant, and 1 if relevant)

$$MAP = \frac{1}{Q} \sum_{i=1}^Q AP_i \quad (2)$$

Where Q is the number of queries and AP_i is the average precision for the i^{th} query.

4.4 Retrieval - MAP Code

```
def calculate_average_precision(relevant_doc_ids,
                                retrieved_doc_ids):
    hit_count = 0
    sum_precisions = 0.0
    for i, doc_id in enumerate(retrieved_doc_ids):
        if doc_id in relevant_doc_ids:
            hit_count += 1
            precision_at_i = hit_count / (i + 1)
            sum_precisions += precision_at_i
        # else: sum_precisions += 0.0
    if len(relevant_doc_ids) == 0:
        return 0.0
    return sum_precisions / len(relevant_doc_ids)
```

5. Qualitative analysis - Information Retrieval

TODO TODO TODO

5.1 Qualitative analysis - Information Retrieval

Problem: Even though there is no relevant information in the document, the system returns some documents

Prompt: "Where can I follow cooking classes"

Output: cinnamon roll glaze taste facts class (0.2410), grandma jayne shrimp dip (0.2369)

5.2 Qualitative analysis - Information Retrieval

Problem: Ignores context of entities in query

Prompt: "How does Gordon Ramsay make his beef Wellington?"

Output: Non of the results were about Ramsay making beef wellington

ID	Score	Words Contained	Words Not Contained
94359	0.2972	ramsay, gordon	beef, wellington
94358	0.2502	ramsay, gordon	beef, wellington
111233	0.2448	beef, wellington	ramsay, gordon
163842	0.2439	beef, wellington	ramsay, gordon
94347	0.2207	ramsay, gordon	beef, wellington
100473	0.2146	beef, wellington	ramsay, gordon
126542	0.2086	beef, wellington	ramsay, gordon
94354	0.2069	ramsay, gordon, beef	wellington
170428	0.2032	ramsay, gordon	beef, wellington
94353	0.2029	ramsay, gordon	beef, wellington

5.3 Qualitative analysis - Information Retrieval

Problem: Can't handle external rare words, like "Paraguay"

Prompt: "Do you know any soups from Paraguay?"

Output: Did not returned any recipe with word "paraguay"

5.4 Qualitative analysis - Information Retrieval

Problem: TF-IDF doesn't handle typos

Prompt: "How do you make **piza**"

Output: Returned single recipe what contains a lot of words "make", but no "pizza"

5.5 Qualitative analysis - Information Retrieval

Problem: Can't capture negation

Prompt: "I do not want to eat pizza, what can I eat instead?"

Output: 36/40 results were about pizza

6. Prompt

Model: mistralai/Mistral-7B-Instruct-v0.2

6.1 Prompt - LLM Instructions - Good

- General context and LLM's goal
- Instructions per type of question
- Response format
- Limitations

See handout for full prompt TODO: add full prompt in handout

6.1 Prompt - LLM Instructions - Bad

Only general context without specific instructions how to answer the question

You are a helpful recipe assistant with access to a database of recipes. The system has already retrieved the most relevant recipes to the user's query using TF-IDF similarity. Your goal is to provide helpful, accurate responses about recipes, cooking techniques, ingredient substitutions, and culinary advice based on the retrieved recipes.

The following recipes have been retrieved as most relevant to the user's query:

{retrieved_recipes}

User Query

{user_query}

6.2 Prompt - Fields used

- name
- description
- ingredients
- steps
- relevance score

6.2 Prompt - Fields used

```
results = retrieve_documents(query, recipes, recipe_ids, k=k)
retrieved_recipes = ""

for idx, (recipe, recipe_id, score) in enumerate(results):
    info=df[df["official_id"] == recipe_id].iloc[0]
    retrieved_recipes+=f"Document {idx}, Score: {score:.4f}\n"
    retrieved_recipes+=f"Name: {info['name']}\n"
    retrieved_recipes+=f"Description:{info['description']}\n"
    retrieved_recipes+=f"Ingredients:{info['ingredients']}\n"
    retrieved_recipes+=f"Steps: {info['steps']}\n\n"

prompt = prompt_template.format(
    retrieved_recipes=retrieved_recipes, user_query=query
)
```


7. Qualitative analysis - LLM

7.1 Qualitative analysis - LLM & RAG

- More or less yes
- But does not stick to rules
- Response are often too general

7.2 Qualitative analysis - LLM Reasoning 1

TODO: add to handout Q: "I'm looking for a spicy pasta sauce, but I forgot the name. I think it was related to Arabs;"

A: "Based on the retrieved recipes, none of them directly mention Arab origin. However, I can help you with a spicy pasta sauce recipe.

One of the recipes in the list, called "spicy pasta sauce," includes ingredients such as spicy sausage, pasta sauce, rotel tomatoes & chilies, minced garlic cloves, and parmesan cheese. To make the sauce, you first cook the spicy sausage in a skillet until no longer pink, then drain it and wipe out the skillet. Return the sausage to the skillet and add the pasta sauce, rotel, and garlic. Simmer for ten minutes and stir in the parmesan cheese. This recipe should result in a spicy pasta sauce with a rich and flavorful taste."

7.2 Qualitative analysis - LLM Reasoning 2

TOOD: add to handout Q: “I want to eat something with cactus. How many recipes do you know?”

A: “Based on the retrieved recipes, I can help you with at least four different recipes that feature cactus as a main ingredient. Here’s a brief summary of each:

...

Each of these recipes offers a unique taste profile, so I recommend you try the one that most appeals to your personal preferences. If you have any specific questions about the recipes or their ingredients, please don’t hesitate to ask.”

7.2 Qualitative analysis - LLM Reasoning 3

TODO: add to handout Q: “Give me some ideas for pasta dishes that contain calamari.”

A: “Based on the retrieved recipes, I suggest trying the following pasta dishes that include calamari:

1. ****Sunday Night Pastaa with Calamari:**** ...
2. ****Calamari Marinara:**** ...

These are our suggestions for pasta dishes with calamari based on the provided recipes. Let me know if you have any further questions or need more details!

If you have dietary restrictions or want to modify the dishes in specific ways, please let me know, and I'll do my best to suggest appropriate adjustments.”

7.2 Qualitative analysis - LLM Reasoning 4

Q: “What temperature should I bake chicken Wellington at?”

A: “The recommended baking temperature for chicken Wellington from the provided recipes is 375 degrees for 25 minutes. However, it’s important to ensure that the internal temperature of the chicken reaches 165°F (74°C) before serving to ensure it is thoroughly cooked. You can use a meat thermometer to check.”

7.2 Qualitative analysis - LLM Reasoning 5

TODO: add to handout Q: “What’s a typical food item from Rhode Island?”

A: “A typical food item from Rhode Island is Coffee Milk. It is a popular comfort food and the official state drink of Rhode Island. To make it, you add coffee-flavored syrup to milk in a large glass, stir well, and taste to add more syrup for a stronger coffee taste, if desired. (Answer taken directly from Document 0.)”

7.3 Qualitative analysis - Hallucination

TODO: add to handout

- Standard prompt, but with no documents provided
"How does Gordon Ramsay make his beef Wellington?"
- Yet LLM answered with recipe claiming they come from our database
- Which suggest that it did not followed rules
- Though at the end it says that it does not have that information in our database about Gordon Ramsay's recipe

Based on the recipes in our database, Beef Wellington is typically made by coating a filet of beef with pate, wrapping it in sliced ham and mushrooms, then covering it with puff pastry. Here are some specific steps:

1. Prepare the beef: ... 7. Serve: ...

As for the specific way Gordon Ramsay makes his Beef Wellington we don't have that information in our database.

However, the method described above is a common way to make this classic dish....

7.4 Qualitative analysis - Score vs No Score

- No particular difference
- For LLM answers look at the handout

8. Neural document embeddings

Model: all-MiniLM-L6-v2

8.1 Neural document embeddings - out of vocabulary - query

Used model's tokenizer to tokenize the query

- Original text: 'kashubian'
- [CLS] ', 'ka', '##shu', '##bian', '[SEP]'

The ## prefix represents subword tokenization, allowing the model to handle words not in its vocabulary.

8.1 Neural document embeddings - out of vocabulary - document

Used model's tokenizer to tokenize the query

- Original text: 'their settlement area is referred to as kashubia they speak the kashubian language which is classified either...'
- '[CLS]', 'their', 'settlement', 'area', 'is', 'referred', 'to', 'as', 'ka', '##shu', '##bia', 'they', 'speak', 'the', 'ka', '##shu', '##bian', 'language', 'which', 'is', 'classified', 'either',


Neural embeddings returned valid results even through the word "kashubian" was not in the vocabulary.

8.2 Neural document embeddings - hyperparams

Grid search over param space and wiki dataset

- `thresholds=np.arange(0.3, 0.60, 0.05),`
- `k_values=np.arange(4, 20, 4),`

Macro F1 Scores for Combinations of k and Threshold



8.2 Neural document embeddings - Metrics - Wiki

K=8, threshold=0.40

- Macro Precision: 0.310
- Macro Recall: 0.352
- Macro F1: 0.260
- Micro Precision: 0.343
- Micro Recall: 0.303
- Micro F1: 0.322
- MAP: 0.216
- Average DCG: 1.566
- Average NDCG: 0.549

8.2 Neural document embeddings - Metrics - Recipes

K=8, threshold=0.40

- Macro Precision: 0.333
- Macro Recall: 0.288
- Macro F1: 0.245
- Micro Precision: 0.366
- Micro Recall: 0.232
- Micro F1: 0.284
- MAP: 0.192
- Average DCG: 1.353
- Average NDCG: 0.565

9. Compression

9.1 Compression - Long Documents

Information Retrieval

- Cover multiple topics
- Contain lots of words
- Limited document length

LLM

- Limited context window
- Needle in a haystack

9.2 Compression - Solutoin

Split documents into chunks

```
from langchain_text_splitters
    import RecursiveCharacterTextSplitter
text_splitter = RecursiveCharacterTextSplitter(
    chunk_size=500,
    chunk_overlap=100,
    length_function=len,
    is_separator_regex=False,
)
```

9. Security

9.1 Security

- Yes, LLM is susceptible TODO...